

Computer Security

A Hands-on Approach

Wenliang Du

Syracuse University

©2017 by Wenliang Du.

All rights reserved. No part of this book may be reproduced, in any form or by any means, without permission in writing from the author.

The author of this book has used his best efforts in preparing this book. These efforts include the development, research, and testing of theories and programs to determine their effectiveness. The author makes no warranty of any kind, expressed or implied, with regard to these programs or the documentation contained in this book. The author shall not be liable in any event for incidental or consequential damages with, or arising out of, the furnishing, performance, or use of these programs.

Printed by: CreateSpace

First Printing: October 2017

ISBN-13: 978-1548367947

ISBN-10: 154836794X

10 9 8 7 6 5 4 3 2 1

Contents

Preface	xiii
About the Author	xvii
Acknowledgments	xix
 I Software Security	 1
1 Set-UID Programs	5
1.1 The Need for Privileged Programs	6
1.1.1 The Password Dilemma	6
1.1.2 Different Types of Privileged Programs	7
1.2 The Set-UID Mechanism	8
1.2.1 A Superman Story	8
1.2.2 How It Works	8
1.2.3 An Example of Set-UID Program	9
1.2.4 How to Ensure Its Security	10
1.2.5 The Set-GID Mechanism	10
1.3 What Can Go Wrong: What Happened to Superman	10
1.4 Attack Surfaces of Set-UID Programs	12
1.4.1 User Inputs: Explicit Inputs	12
1.4.2 System Inputs	13
1.4.3 Environment Variables: Hidden Inputs	13
1.4.4 Capability Leaking	14
1.5 Invoking Other Programs	16
1.5.1 Unsafe Approach: Using <code>system()</code>	16
1.5.2 Safe Approach: Using <code>execve()</code>	18
1.5.3 Invoking External Commands in Other Languages	19
1.5.4 Lessons Learned: Principle of Isolation	20
1.6 Principle of Least Privilege	20
1.7 Summary	21
 2 Environment Variables and Attacks	 23
2.1 Environment Variables	24
2.1.1 How to Access Environment Variables	24
2.1.2 How a Process Gets Its Environment Variables	25

2.1.3	Memory Location for Environment Variables	26
2.1.4	Shell Variables and Environment Variables	27
2.2	Attack Surface	30
2.3	Attacks via Dynamic Linker	32
2.3.1	Static and Dynamic Linking	32
2.3.2	Case Study: LD_PRELOAD and LD_LIBRARY_PATH	34
2.3.3	Case Study: OS X Dynamic Linker	36
2.4	Attack via External Program	37
2.4.1	Two Typical Ways to Invoke External Programs	37
2.4.2	Case Study: the PATH environment variable	37
2.4.3	Reduce Attack Surface	38
2.5	Attack via Library	38
2.5.1	Case Study - Locale in UNIX	39
2.6	Application Code	39
2.6.1	Case Study - Using <code>getenv()</code> in Application Code	40
2.7	Set-UID Approach versus Service Approach	41
2.8	Summary	42
3	Shellshock Attack	43
3.1	Background: Shell Functions	44
3.2	The Shellshock Vulnerability	46
3.2.1	The Shellshock Bug	46
3.2.2	Mistake in the Bash Source Code	46
3.2.3	Exploiting the Shellshock vulnerability	47
3.3	Shellshock Attack on Set-UID Programs	48
3.4	Shellshock Attack on CGI Programs	49
3.4.1	Experiment Environment Setup	50
3.4.2	How Web Server Invokes CGI Programs	50
3.4.3	How Attacker Sends Data to Bash	51
3.4.4	Launching the Shellshock Attack	52
3.4.5	Creating Reverse Shell	53
3.5	Remote Attack on PHP	55
3.6	Summary	56
4	Buffer Overflow Attack	57
4.1	Program Memory Layout	58
4.2	Stack and Function Invocation	59
4.2.1	Stack Memory Layout	60
4.2.2	Frame Pointer	60
4.3	Stack Buffer-Overflow Attack	62
4.3.1	Copy Data to Buffer	62
4.3.2	Buffer Overflow	63
4.3.3	Exploiting a Buffer Overflow Vulnerability	64
4.4	Setup for Our Experiment	66
4.4.1	Disable Address Randomization	66
4.4.2	Vulnerable Program	66
4.5	Conduct Buffer-Overflow Attack	67
4.5.1	Finding the Address of the Injected Code	68

4.5.2	Improving Chances of Guessing	69
4.5.3	Finding the Address Without Guessing	69
4.5.4	Constructing the Input File	71
4.6	Writing a Shellcode	73
4.6.1	Writing Malicious Code Using C	73
4.6.2	Writing a Shellcode: Main Idea	74
4.6.3	Explanation of a Shellcode Example	74
4.7	Countermeasures: Overview	77
4.8	Address Randomization	79
4.8.1	Address Randomization on Linux	79
4.8.2	Effectiveness of Address Randomization	81
4.9	StackGuard	82
4.9.1	The Observation and the Idea	82
4.9.2	Manually Adding Code to Function	83
4.9.3	StackGuard Implementation in gcc	84
4.10	Summary	87
5	Return-to-libc Attack	89
5.1	Introduction	90
5.2	The Attack Experiment: Setup	91
5.3	Launch the Return-to-libc Attack: Part I	93
5.3.1	Task A: Find the Address of the <code>system()</code> Function	93
5.3.2	Task B: Find the Address of the String <code>"/bin/sh"</code>	94
5.4	Launch the Return-to-libc Attack: Part II	95
5.4.1	Function Prologue	95
5.4.2	Function Epilogue	96
5.4.3	Function Prologue and Epilogue Example	97
5.4.4	Perform Task C	98
5.4.5	Construct Malicious Input	99
5.4.6	Launch the Attack	100
5.5	Summary	101
6	Format String Vulnerability	103
6.1	Functions with Variable Number of Arguments	104
6.1.1	How to Access Optional Arguments	104
6.1.2	How <code>printf()</code> Accesses Optional Arguments	106
6.2	Format String with Missing Optional Arguments	107
6.3	Vulnerable Program and Experiment Setup	109
6.4	Exploiting the Format String Vulnerability	110
6.4.1	Attack 1: Crash Program	110
6.4.2	Attack 2: Print out Data on the Stack	111
6.4.3	Attack 3: Change the Program's Data in the Memory	111
6.4.4	Attack 4: Change the Program's Data to a Specific Value	113
6.4.5	Attack 4 (Continuation): A Much Faster Approach	114
6.4.6	Attack 5: Inject Malicious Code	116
6.4.7	Reducing the Size of Format String	118
6.5	Countermeasures	119
6.5.1	Developer	119

6.5.2	Compiler	119
6.5.3	Address Randomization	120
6.6	Summary	120
7	Race Condition Vulnerability	123
7.1	The General Race Condition Problem	124
7.2	Race Condition Vulnerability	125
7.3	Experiment Setup	127
7.4	Exploiting Race Condition Vulnerabilities	128
7.4.1	Choose a Target File	128
7.4.2	Launch Attack	129
7.4.3	Monitor the Result	130
7.4.4	Running the Exploit	130
7.5	Countermeasures	131
7.5.1	Atomic Operation	131
7.5.2	Repeating Check and Use	132
7.5.3	Sticky Symlink Protection	133
7.5.4	Principle of Least Privilege	135
7.6	Summary	136
8	Dirty COW	137
8.1	Memory Mapping using <code>mmap()</code>	138
8.2	<code>MAP_SHARED</code> , <code>MAP_PRIVATE</code> and Copy On Write	139
8.3	Discard the Copied Memory	141
8.4	Mapping Read-Only Files	141
8.5	The Dirty COW Vulnerability	143
8.6	Exploiting the Dirty COW Vulnerability	144
8.6.1	Selecting <code>/etc/passwd</code> as Target File	145
8.6.2	Set Up the Memory Mapping and Threads	145
8.6.3	The <code>write</code> Thread	146
8.6.4	The <code>madvise</code> Thread	147
8.6.5	The Attack Result	147
8.7	Summary	148
II	Web Security	149
9	Cross Site Request Forgery	153
9.1	Cross-Site Requests and Its Problems	154
9.2	Cross-Site Request Forgery Attack	155
9.3	CSRF Attacks on HTTP GET Services	156
9.3.1	HTTP GET and POST Services	156
9.3.2	The Basic Idea of CSRF Attacks	157
9.3.3	Attack on Elgg's Add-friend Service	157
9.4	CSRF Attacks on HTTP POST Services	159
9.4.1	Constructing a POST Request Using JavaScript	159
9.4.2	Attack on Elgg's Edit-Profile Service	160
9.5	Countermeasures	162

9.5.1	Using the <code>referer</code> Header	163
9.5.2	Same-Site Cookies	163
9.5.3	Secret Token	163
9.5.4	Case Study: Elgg's Countermeasures	164
9.6	Summary	166
10	Cross-Site Scripting Attack	167
10.1	The Cross-Site Scripting Attack	168
10.1.1	Non-persistent (Reflected) XSS Attack	168
10.1.2	Persistent XSS Attack	170
10.1.3	What damage can XSS cause?	170
10.2	XSS Attacks in Action	171
10.2.1	Prelude: Injecting JavaScript Code	171
10.2.2	Use XSS Attacks to Befriend with Others	172
10.2.3	Use XSS Attacks to Change Other People's Profiles	175
10.3	Achieving Self-Propagation	177
10.3.1	Creating a Self-Propagating XSS Worm: the DOM Approach	178
10.3.2	Create a Self-Propagating Worm: the Link Approach	180
10.4	Preventing XSS attacks	181
10.5	Summary	182
11	SQL Injection Attack	183
11.1	A Brief Tutorial of SQL	184
11.1.1	Log in to MySQL	184
11.1.2	Create a Database	184
11.1.3	CREATE a Table	184
11.1.4	INSERT a Row	185
11.1.5	The SELECT Statement	185
11.1.6	WHERE Clause	186
11.1.7	UPDATE SQL Statement	187
11.1.8	Comments in SQL Statements	187
11.2	Interacting with Database in Web Application	188
11.2.1	Getting Data from User	188
11.2.2	Getting Data From Database	189
11.3	Launching SQL Injection Attacks	191
11.3.1	Attack Using <code>cURL</code>	192
11.3.2	Modify Database	192
11.3.3	Multiple SQL Statements	193
11.4	The Fundamental Cause	194
11.5	Countermeasures	197
11.5.1	Filtering and Encoding Data	197
11.5.2	Prepared Statement	197
11.6	Summary	199

III	Network Security	201
12	Packet Sniffing and Spoofing	205
12.1	How Packets Are Received	206
12.1.1	Network Interface Card (NIC)	206
12.1.2	BSD Packet Filter (BPF)	207
12.2	Packet Sniffing	208
12.2.1	Receiving Packets Using Sockets	208
12.2.2	Packet Sniffing using Raw Sockets	209
12.2.3	Packet Sniffing Using the <code>pcap</code> API	211
12.2.4	Processing Captured Packet	212
12.3	Packet Spoofing	215
12.3.1	Sending Normal Packets Using Socket	215
12.3.2	Sending Spoofed Packets Using Raw Sockets	216
12.3.3	Constructing ICMP Packets	218
12.3.4	Constructing UDP Packets	219
12.4	Snoofing: Sniffing and Spoofing	221
12.5	Endianness	223
12.6	Calculating Checksum	224
12.7	Summary	226
13	Attacks on the TCP Protocol	227
13.1	How the TCP Protocol Works	228
13.1.1	TCP Client Program	228
13.1.2	TCP Server Program	229
13.1.3	Data Transmission: Under the Hood	232
13.1.4	TCP Header	233
13.2	SYN Flooding Attack	234
13.2.1	TCP Three-Way Handshake Protocol	234
13.2.2	The SYN Flooding Attack	235
13.2.3	Launching the SYN Flooding Attack	236
13.2.4	Launching SYN Flooding Attacks Using Our Own Code	238
13.2.5	Countermeasure	239
13.3	TCP Reset Attack	240
13.3.1	Closing TCP Connections	240
13.3.2	How the Attack Works	241
13.3.3	Launching the TCP Reset Attack: Setup	241
13.3.4	TCP Reset Attack on <code>Telnet</code> connections	242
13.3.5	TCP Reset Attack on SSH connections	244
13.3.6	TCP Reset Attack on Video-Streaming Connections	244
13.4	TCP Session Hijacking Attack	245
13.4.1	TCP Session and Session Hijacking	246
13.4.2	Launching the TCP Session Hijacking Attack	247
13.4.3	What Happens to the Hijacked TCP Connection	249
13.4.4	Causing More Damage	251
13.4.5	Creating Reverse Shell	251
13.5	Summary	252

14 Firewall	255
14.1 Introduction	256
14.2 Types of Firewalls	257
14.2.1 Packet Filter	257
14.2.2 Stateful Firewall	258
14.2.3 Application/Proxy Firewall	258
14.3 Building a Simple Firewall using Netfilter	258
14.3.1 Writing Loadable Kernel Modules	259
14.3.2 Compiling Kernel Modules	260
14.3.3 Installing Kernel Modules	260
14.4 Netfilter	261
14.4.1 netfilter Hooks for IPv4	262
14.4.2 Implementing a Simple Packet Filter Firewall	262
14.5 The iptables Firewall in Linux	265
14.5.1 The structure of the iptables Firewall	265
14.5.2 Traversing Chains and Rule Matching	266
14.5.3 iptables Extensions	267
14.5.4 Building a Simple Firewall	268
14.6 Stateful Firewall using Connection Tracking	270
14.6.1 Stateful Firewall	270
14.6.2 The Connection Tracking Framework in Linux	271
14.6.3 Example: Set up a Stateful Firewall	272
14.7 Application/Proxy Firewall and Web Proxy	272
14.8 Evading Firewalls	273
14.8.1 Using SSH Tunneling to Evade Firewalls	273
14.8.2 Dynamic Port Forwarding	275
14.8.3 Using VPN to Evade Firewall	276
14.9 Summary	276
15 Domain Name System (DNS) and Attacks	279
15.1 DNS Hierarchy, Zones, and Servers	280
15.1.1 DNS Domain Hierarchy	280
15.1.2 DNS Zone	281
15.1.3 Authoritative Name Servers	282
15.1.4 The Organization of Zones on the Internet	282
15.2 DNS Query Process	283
15.2.1 Local DNS Files	284
15.2.2 Local DNS Server and the Iterative Query Process	285
15.3 Set Up DNS Server and Experiment Environment	287
15.3.1 Configure the User Machine	287
15.3.2 Configure the Local DNS server	289
15.3.3 Set Up Zones in the Local DNS Server	290
15.4 DNS Attacks: Overview	292
15.5 Local DNS Cache Poisoning Attack	294
15.6 Remote DNS Cache Poisoning Attack	296
15.6.1 The Kaminsky Attack	297
15.6.2 Construct the IP and UDP headers of DNS reply	299
15.6.3 Construct the DNS Header and Payload	300

15.6.4	Result Verification	301
15.7	Reply Forgery Attacks from Malicious DNS Servers	302
15.7.1	Fake Data in the Additional Section	302
15.7.2	Fake Data in the Authority Section	302
15.7.3	Using Both Sections	303
15.7.4	Fake Data in the Answer Section: for Reverse DNS Lookup	304
15.8	Protection Against DNS Cache Poisoning Attacks	306
15.8.1	DNSSEC	306
15.8.2	TLS/SSL Solution	306
15.9	Denial of Service Attacks on DNS Servers	308
15.9.1	Attacks on the Root and TLD Servers	308
15.9.2	Attacks on Nameservers of a Particular Domain	309
15.10	Summary	310
16	Virtual Private Network	311
16.1	Introduction	312
16.1.1	Virtual Private Network	312
16.1.2	How a Virtual Private Network Works	314
16.2	An Overview of How TLS/SSL VPN Works	315
16.2.1	Establishing A TLS/SSL Tunnel	316
16.2.2	Forwarding IP packets	316
16.2.3	Releasing IP Packets	317
16.3	How TLS/SSL VPN Works: Details	318
16.3.1	Virtual Network Interfaces	318
16.3.2	Creating a TUN Interface	319
16.3.3	Routing Packets to a TUN Interface	321
16.3.4	Reading and Writing Operations on the TUN Interface	322
16.3.5	Forwarding Packets via the Tunnel	322
16.3.6	Packet's Return Trip	323
16.4	Building a VPN	323
16.4.1	Establish the Tunnel	324
16.4.2	Monitoring File Descriptors	326
16.4.3	From TUN To Tunnel	326
16.4.4	From Tunnel to TUN	327
16.4.5	Bring Everything Together	327
16.5	Setting Up a VPN	328
16.5.1	Network Configuration	328
16.5.2	Testing VPN	330
16.6	Using VPN to Bypass Egress Firewall	332
16.6.1	Network Setup	332
16.6.2	Setting Up VPN to Bypass Firewall	334
16.7	Summary	335
17	The Heartbleed Bug and Attack	337
17.1	Background: the Heartbeat Protocol	338
17.2	Launch the Heartbleed Attack	340
17.2.1	Attack Environment and Setup	340
17.2.2	Launch an Attack	341

17.3	Fixing the Heartbleed Bug	343
17.4	Summary	343
18	Public Key Infrastructure	345
18.1	Attack on Public Key Cryptography	346
18.1.1	Man-in-the-Middle (MITM) Attack	346
18.1.2	Defeating MITM Attacks	347
18.1.3	Public Key Infrastructure	347
18.2	Public Key Certificates	348
18.2.1	X.509 Digital Certificate	348
18.2.2	Get Certificate from a Real Server	349
18.3	Certificate Authority (CA)	350
18.3.1	Being a CA	351
18.3.2	Getting X.509 Certificate from CA	352
18.3.3	Deploying Public Key Certificate in Web Server	354
18.3.4	Apache Setup for HTTPS	356
18.4	Root and Intermediate Certificate Authorities	356
18.4.1	Root CAs and Self-Signed Certificate	356
18.4.2	Intermediate CAs and Chain of Trust	357
18.4.3	Creating Certificates for Intermediate CA	358
18.4.4	Apache Setup	359
18.4.5	Trusted CAs in the Real World	359
18.5	How PKI Defeats the MITM Attack	360
18.5.1	Attacker Forwards the Authentic Certificate	360
18.5.2	Attacker Creates a Fake Certificate	361
18.5.3	Attacker Sends Its Own Certificate	361
18.5.4	The Man-In-The-Middle Proxy	362
18.6	Attacks on the Public-Key Infrastructure	363
18.6.1	Attack on CA's Verification Process	364
18.6.2	Attack on CA's Signing Process	365
18.6.3	Attacks on the Algorithms	365
18.6.4	Attacks on User Confirmation	366
18.7	Types of Digital Certificates	367
18.7.1	Domain Validated Certificates (DV)	367
18.7.2	Organizational Validated Certificates (OV)	368
18.7.3	Extended Validated Certificates (EV)	368
18.8	Summary	369
19	Transport Layer Security	371
19.1	Overview of TLS	372
19.2	TLS Handshake	373
19.2.1	Overview of the TLS Handshake Protocol	373
19.2.2	Certificate Verification	375
19.2.3	Key Generation and Exchange	375
19.3	TLS Data Transmission	377
19.3.1	Sending Data with TLS Record Protocol	377
19.3.2	Receiving Data with TLS Record Protocol	378
19.4	TLS Programming: A Client Program	379

19.4.1	The Overall Picture	379
19.4.2	TLS Initialization	380
19.4.3	TCP Connection Setup	381
19.4.4	TLS Handshake	382
19.4.5	Application Data Transmission	383
19.4.6	Set Up the Certificate Folder	383
19.4.7	The Complete Client Code	385
19.5	Verifying Server's Hostname	386
19.5.1	Modified Client Code	386
19.5.2	An Experiment: Man-In-The-Middle Attack	387
19.5.3	Hostname Checking	389
19.6	TLS Programming: the Server Side	390
19.6.1	TLS Setup	390
19.6.2	TCP Setup	392
19.6.3	TLS Handshake	392
19.6.4	TLS Data Transmission	393
19.6.5	Testing	394
19.7	Summary	394

Preface

This book is for students, computer scientists, computer engineers, programmers, software developers, network and system administrators, and others who want to learn the principles of computer security and understand how various security attacks and countermeasures work. Equipped with the knowledge from this book, readers will be able to design and implement software systems and applications that are secure against attacks. They will also be able to evaluate the risks faced by computer and network systems, detect common vulnerabilities in software, use proper methods to protect their systems and networks, and more importantly, apply the learned security principles to solve real-world problems.

The author strongly believes in “learning by doing”, so the book takes a hands-on approach. For each security principle, the book uses a series of hands-on activities to help explain the principle; readers can *touch*, *play with*, and *experiment with* the principle, instead of just reading about it. For instance, if a security principle involves an attack, the book guides readers to actually launch the attack (in a contained environment). If a principle involves a security mechanism, such as firewall or Virtual Private Network (VPN), the book guides readers to implement a mini-firewall or mini-VPN. Readers can learn better from such hands-on activities.

All the hands-on activities are conducted in a virtual machine image provided by the author. They can be downloaded from this URL: <http://www.cis.syr.edu/~wedu/seed/>. Everything needed for the activities have already been set up; readers just need to download the VM (free), launch it using `VirtualBox`, and they can immediately work on the activities covered in the book. This book is based on the `Ubuntu12.04` VM image. The author will regularly upgrade the VM image in every few years.

Most of the activities in the book are based on the author’s SEED labs, which are widely used by instructors all over the world. These labs are the results of 15 years’ research, development, and testing efforts conducted by the author and his students in a project called SEED, which has been funded by the National Science Foundation since 2002.

The Organization of the Book

The book are organized in three broad topics: software security, web security, and network security. Software security and web security cover some of the well-known vulnerabilities and attacks in general software and web applications, including a few recent attacks, such as the Shellshock and Dirty COW attacks. By learning these topics, readers can understand why a computer or a program can be attacked, what is under the hood in these attacks, and how to

write better programs so they are immune or more resilient to attacks. The network security part focuses on the security principles related to the Internet. It not only covers some of well-known attacks on the Internet, but also covers important defense mechanisms, such as firewall, VPN, and PKI.

The book is not intended to cover every attack or security measure. The topics covered in the book are representative in terms of covering the fundamental security principles. Some of the topics, such as cryptography, system security, and mobile security, are left out for the time being, so the publication of this book will not be delayed for another one or two years. Some of these topics will be added in future editions. The contents of this book are sufficient for the courses that cover the fundamental principles of cybersecurity. For example, two of the author's courses (*Computer Security* and *Internet Security*) are based on the contents of this book. These two courses are taught at both undergraduate and graduate levels.

While some chapters depend on previous chapters, most chapters are self-contained, and can be read independently. The following list describes the partial dependence relationship among chapters.

- Chapter 1 (Set-UID Programs) is the basis for most chapters in software security. This chapter describes how the Set-UID mechanism works and gives an overview of the attacks that can be launched against this type of privileged program. Although there are many other types of privileged program, we use this type of program to explain how various attacks work.
- Chapter 2 (Environment Variables) is the basis for Chapter 3 (Shellshock).
- Chapter 4 (Buffer Overflow) is the basis for Chapter 5 (Return-to-libc Attack), because return-to-libc attacks defeat one of the countermeasures covered in Chapter 4.
- Chapter 7 (Race Condition) and Chapter 8 (Dirty COW) are both related to the race condition vulnerability, but we suggest readers to read Chapter 7 first, as it is easier to understand.
- Chapter 12 (Sniffing and Spoofing) is the basis for most of the network attacks covered in the book, so it should be read first before the other chapters in Network Security.
- Chapter 18 (Public Key Infrastructure) is the basis for Chapter 19 (Transport Layer Protocol).

The History of the SEED labs

“I hear and I forget. I see and I remember. I do and I understand”. This famous saying, by Chinese philosophy Confucius (551 BC – 479 BC), has been a motto for many educators, who firmly believe that learning must be grounded in experience. This is particularly true for computer security education. Sixteen years ago, with this motto taken to the heart, and a desire to become an excellent instructor in computer security, The author searched the Web, looking for hands-on projects that he could use for his security classes. He could only find a few, but they came from various places, and were incoherent; their coverage of security topics was quit narrow, even jointly, and the lab environments they used were not easy nor inexpensive to set up.

Determined, he decided to develop his own hands-on exercises (called labs in short), not one lab, but many of them, covering a wide spectrum of security topics; not just for his own

use, but for many other instructors who share the same teaching philosophy as he does. All the labs should be based on one unified environment, so students do not need to spend too much time learning a new environment for different labs. Moreover, the lab environment should be easy and inexpensive to set up, so instructors are not hindered even if they have limited time or resources.

With the above goals in mind and an initial grant from NSF (\$74,984.00, Award No. 0231122), he started the journey in 2002, naming the project as SEED (standing for SEcurity EDucation). Ten years later, after another NSF grant (\$451,682, Award No. 0618680) and the help from over 20 students, he has developed about 30 SEED labs, covering many security topics, including vulnerabilities, attacks, software security, system security, network security, web security, access control, cryptography, mobile security, etc. Most SEED labs have gone through multiple development-trial cycles—development, trial, improvement, and trial again—in actual courses at Syracuse University and many other institutes.

The SEED project has been quite successful. As of now, more than 600 instructors worldwide told the author that they have used some of the SEED labs; more people simply used the SEED labs without telling (which is perfectly fine), as all the SEED lab materials and the lab environment are available online, free of charge. To help others use the SEED labs, NSF gave the author another grant (\$863,385.00, Award No. 1303306), so he can organize two training workshops each year and fund those who are interested to come to attend the workshops. Every year, about 70 instructors attended the workshops.

About the Author



Wenliang (Kevin) Du, PhD, received his bachelor's degree from the University of Science and Technology of China in 1993. After getting a Master's degree from Florida International University, he attended Purdue University from 1996 to 2001, and received his PhD degree in computer science. He became an assistant professor at Syracuse University after the graduation. He is currently a full professor in the Department of Electrical Engineering and Computer Science.

Professor Du has taught courses in cybersecurity at both undergraduate and graduate levels since 2001. As a firm believer of “learning by doing”, he has developed over 30 hands-on labs called SEED labs, so students can gain first-hand experiences on security attacks, countermeasures, and fundamental security principles. These labs are now widely known; more than six hundred universities, colleges, and high schools worldwide are using or have used these labs. In 2010, the SEED project was highlighted by the National Science Foundation in a report sent to the Congress. The report, titled “New Challenges, New Strategies: Building Excellence in Undergraduate STEM Education (Page 16)”, highlights “17 projects that represent cutting-edge creativity in undergraduate STEM classes nationwide”. Due to the impact of the SEED labs, he was given the “2017 Academic Leadership” award from the *21st Colloquium for Information System Security Education*.

Professor Du works in the area of computer and network security, with specific interests in system security. He has published over 100 technical papers. As of August 2017, his research work has been cited for over 12,500 times (based on Google Scholar). He is a recipient of the ACM CCS Test-of-Time Award in 2013 due to the impact of one of his papers published in 2003. His current research focuses on smartphone security. He has identified a number of security problems in the design and implementation of the Android operating system. He also developed novel mechanisms to enhance the system security of smartphones.

Acknowledgments

I would like to thank the National Science Foundation for providing the funding support for my SEED project, which laid the foundation for this book. Since 2002, three NSF grants supported the SEED project, including Award No. 0231122, 0618680, and 1303306. I especially thank the Program Director Dr. Victor P. Piotrowski for his leadership in cybsecurity education and for putting the trust in my SEED project.

The SEED project is built on the joint effort of many of my students over the past 15 years. I would like to acknowledge the following students for their contributions: Dr. Yousra Aafer, Amit Ahlawat, Francis Akowuah, Swapnil Bhalode, Ashok Bommisetti, Sudheer Bysani, Bandan Das, Nishant Doshi, Jinkai Gao, Hao Hao, Lin Huang, Sridhar Iyer, Apoorva Iyer, Dr. Karthick Jayaraman, Yuexin (Eric) Jiang, Xing Jin, Vishtasp Jokhi, Sharath B. Koratikere, Dr. Tongbo Luo, Sankara Narayanan, Nagesh Gautam Peri, Karankumar H. Patel, Amey Patil, Balamurugan Rajagopalan, Dr. Paul Ratazzi, Divyakaran Sachar, Mingdong Shang, Sunil Vajir, Dr. Ronghua Wang, Shaonan Wang, Yifei Wang, Zhenyu Wang, Kailiang Ying, Haichao Zhang, Dr. Xiao Zhang, Zhuo Zhang, and Dr. Zutao Zhu.

I would like to acknowledge all the instructors who have used my SEED labs in their classes, as well as those who attended my workshops. Many of them send me encouraging words, suggestions, and feedbacks; they also helped spread the words about my SEED labs. They made my work meaningful, and inspired me to keep moving forward in my project.

Most importantly, I would like to thank my family for their support, for their trust in me, and for the sacrifice of family time due to the writing of this book.